

Script: To run or view read “README.md” or “README.pdf” in “ECE 3195 Project Code and data.zip”

Data:

Computed data (model, accuracy, loss, etc.) is in for of tensorboard and can be found in “tensorboard” folder in “ECE 3195 Project Code and data.zip”.

View is using “`tensorboard -port <port> --logdir tensorboard`” command.

Analysis:

- The model from Part 1 got an accuracy of 72% which is impressive for a small model when compared to VGG11 from part 2, which shows how important it is to tune hyperparameter correctly to get good accuracy.
- In part 2 sub-part 1 (a), VGG11 with no hyper-parameter tuning got an accuracy of 75% even after 100 epoch which is comparable to the model from part 1, again showing the importance of tuning of hyper-parameter.
- Part 2 sub-part 1 (b) and Part 2 sub-part 2 with BatchNorm, BatchNorm decreases the variation of accuracy and loss between each step hence making the training more stable. This is done by normalization of layers’ inputs over the batch.
- Part 2 sub-part 1 (c), using LeakyReLU instead of ReLU, LeakyReLU decreases the number of epochs needed to train by giving a small negative activation when wrong inputs are sent, which helps it to learn better and faster during the training phase.
- Part 2 sub-part 1 (a), Change the optimizer, learning rate, and other optimizer parameters also greatly impact the learning process and final accuracy and generality of the model.
- Part 2 sub-part 1 (b), without BatchNorm changing batch size, doesn’t change anything, but with BatchNorm the change can be noticeable (larger batch size may result in a more accurate model).
- Part 2 sub-part 1 (c), how the model is initialized also plays an important role, since deep learning occurs in a domain with multiple local minima, so with different initialization may result in the trained model being in different minima.
- Part 2 sub-part 1 (d), Dropout layers help in the distribution of learned parameters or information thought out the layer and remove the possibility of the layer with only a few significant neurons with an important or significant role. It also helps the model generalize better resulting in better testing accuracy. This is done by randomly dropping neurons with some given probability.

Results:

Part 1: Basis PyTorch

Parameters :

Name: "assignment_p1"

Epoch = 14

Batch Size = 64

Loss function: NLLL Loss

Optimizer: Adadelata

Model:

```
NNModel(  
  (conv_nn): Sequential(  
    (0): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1))  
    (1): ReLU(inplace=True)  
    (2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))  
    (3): ReLU(inplace=True)  
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0)  
    (5): Dropout(p=0.25, inplace=False)  
  )  
  (flatten): Sequential(  
    (0): Flatten(start_dim=1, end_dim=-1)  
  )  
  (fc_nn): Sequential(  
    (0): Linear(in_features=12544, out_features=128, bias=True)  
    (1): ReLU(inplace=True)  
    (2): Dropout(p=0.5, inplace=False)  
    (3): Linear(in_features=128, out_features=10, bias=True)  
    (4): LogSoftmax(dim=1)  
  )  
)
```




Figure 1. Testing Accuracy (72%)




Figure 2. Training Accuracy (73%)




Figure 3. Testing Loss




Figure 4. Training Loss

Part 2, Sub-Part 0:

Parameters :

Name: "assignment_p2_0"

Epoch = 14

Batch Size = 64

Loss function: Cross Entropy Loss

Optimizer: Adam (betas:(0.5, 0.999), lr: 2e-4)

Initialization function: kaiming_normal_

Model:

```
VGG11Model(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): ReLU(inplace=True)  
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): ReLU(inplace=True)  
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (6): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): ReLU(inplace=True)  
    (8): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): ReLU(inplace=True)  
    (10): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (11): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (12): ReLU(inplace=True)  
    (13): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (14): ReLU(inplace=True)  
    (15): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (16): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (17): ReLU(inplace=True)  
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (19): ReLU(inplace=True)  
    (20): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avg_pool): AdaptiveAvgPool2d(output_size=(7, 7))  
  (classifier): Sequential(  
    (0): Linear(in_features=25088, out_features=4096, bias=True)  
    (1): ReLU(inplace=True)  
    (2): Dropout(p=0.5, inplace=False)  
    (3): Linear(in_features=4096, out_features=4096, bias=True)  
    (4): ReLU(inplace=True)  
    (5): Dropout(p=0.5, inplace=False)  
    (6): Linear(in_features=4096, out_features=10, bias=True)  
  )  
)
```




Figure 8. Testing Accuracy (74%)




Figure 7. Training Accuracy (97%)




Figure 6. Testing Loss




Figure 5. Training Loss

Part 2, Sub-Part 1 (a):

Parameters :

Name: "assignment_p2_1a"

Epoch = 100

Batch Size = 64

Loss function: Cross Entropy Loss

Optimizer: Adam (betas:(0.5, 0.999), lr: 2e-4)

Initialization function: kaiming_normal_

Model:

```
VGG11Model(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): ReLU(inplace=True)  
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): ReLU(inplace=True)  
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (6): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): ReLU(inplace=True)  
    (8): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): ReLU(inplace=True)  
    (10): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (11): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (12): ReLU(inplace=True)  
    (13): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (14): ReLU(inplace=True)  
    (15): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (16): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (17): ReLU(inplace=True)  
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (19): ReLU(inplace=True)  
    (20): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avg_pool): AdaptiveAvgPool2d(output_size=(7, 7))  
  (classifier): Sequential(  
    (0): Linear(in_features=25088, out_features=4096, bias=True)  
    (1): ReLU(inplace=True)  
    (2): Dropout(p=0.5, inplace=False)  
    (3): Linear(in_features=4096, out_features=4096, bias=True)  
    (4): ReLU(inplace=True)  
    (5): Dropout(p=0.5, inplace=False)  
    (6): Linear(in_features=4096, out_features=10, bias=True)  
  )  
)
```




Figure 9. Testing Accuracy (76.63%)




Figure 12. Training Accuracy (99.14%)




Figure 11. Testing Loss




Figure 10. Training Loss

Part 2, Sub-Part 1 (b):

Parameters :

Name: "assignment_p2_1b"

Epoch = 100

Batch Size = 64

Loss function: Cross Entropy Loss

Optimizer: Adam (betas:(0.5, 0.999), lr: 2e-4)

Initialization function: kaiming_normal_

Model:

```
VGG11Model(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU(inplace=True)  
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (4): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (5): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (6): ReLU(inplace=True)  
    (7): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (8): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (10): ReLU(inplace=True)  
    (11): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (12): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (13): ReLU(inplace=True)  
    (14): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (15): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (16): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (17): ReLU(inplace=True)  
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (19): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (20): ReLU(inplace=True)  
    (21): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (22): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (23): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (24): ReLU(inplace=True)  
    (25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (26): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (27): ReLU(inplace=True)  
    (28): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avg_pool): AdaptiveAvgPool2d(output_size=(7, 7))  
  (classifier): Sequential(  
    (0): Linear(in_features=25088, out_features=4096, bias=True)  
    (1): ReLU(inplace=True)  
    (2): Dropout(p=0.5, inplace=False)  
    (3): Linear(in_features=4096, out_features=4096, bias=True)  
    (4): ReLU(inplace=True)  
    (5): Dropout(p=0.5, inplace=False)  
    (6): Linear(in_features=4096, out_features=10, bias=True)  
  )  
)
```




Figure 13. Testing Accuracy (80.92%)




Figure 16. Training Accuracy (99.41%)




Figure 15. Testing Loss




Figure 14. Training Loss

Part 2, Sub-Part 1 (c):

Parameters :

Name: "assignment_p2_1c"

Epoch = 100

Batch Size = 64

Loss function: Cross Entropy Loss

Optimizer: Adam (betas:(0.5, 0.999), lr: 2e-4)

Initialization function: kaiming_normal_

Model:

```
VGG11Model(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): LeakyReLU(negative_slope=0.01, inplace=True)  
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (6): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): LeakyReLU(negative_slope=0.01, inplace=True)  
    (8): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): LeakyReLU(negative_slope=0.01, inplace=True)  
    (10): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (11): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (12): LeakyReLU(negative_slope=0.01, inplace=True)  
    (13): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (14): LeakyReLU(negative_slope=0.01, inplace=True)  
    (15): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (16): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (17): LeakyReLU(negative_slope=0.01, inplace=True)  
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (19): LeakyReLU(negative_slope=0.01, inplace=True)  
    (20): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avg_pool): AdaptiveAvgPool2d(output_size=(7, 7))  
  (classifier): Sequential(  
    (0): Linear(in_features=25088, out_features=4096, bias=True)  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): Dropout(p=0.5, inplace=False)  
    (3): Linear(in_features=4096, out_features=4096, bias=True)  
    (4): LeakyReLU(negative_slope=0.01, inplace=True)  
    (5): Dropout(p=0.5, inplace=False)  
    (6): Linear(in_features=4096, out_features=10, bias=True)  
  )  
)
```




Figure 17. Testing Accuracy (77.02%)




Figure 20. Training Accuracy (98.51%)




Figure 19. Testing Loss




Figure 18. Training Loss

Part 2, Sub-Part 1 (c) with BatchNorm:

Parameters :

Name: "assignment_p2_1c_b"

Epoch = 100

Batch Size = 64

Loss function: Cross Entropy Loss

Optimizer: Adam (betas:(0.5, 0.999), lr: 2e-4)

Initialization function: kaiming_normal_

Model:

```
VGG11Model(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): LeakyReLU(negative_slope=0.01, inplace=True)  
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (4): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (5): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (6): LeakyReLU(negative_slope=0.01, inplace=True)  
    (7): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (8): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (10): LeakyReLU(negative_slope=0.01, inplace=True)  
    (11): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (12): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (13): LeakyReLU(negative_slope=0.01, inplace=True)  
    (14): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (15): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (16): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (17): LeakyReLU(negative_slope=0.01, inplace=True)  
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (19): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (20): LeakyReLU(negative_slope=0.01, inplace=True)  
    (21): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (22): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (23): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (24): LeakyReLU(negative_slope=0.01, inplace=True)  
    (25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (26): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (27): LeakyReLU(negative_slope=0.01, inplace=True)  
    (28): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avg_pool): AdaptiveAvgPool2d(output_size=(7, 7))  
  (classifier): Sequential(  
    (0): Linear(in_features=25088, out_features=4096, bias=True)  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): Dropout(p=0.5, inplace=False)  
    (3): Linear(in_features=4096, out_features=4096, bias=True)  
    (4): LeakyReLU(negative_slope=0.01, inplace=True)  
    (5): Dropout(p=0.5, inplace=False)  
    (6): Linear(in_features=4096, out_features=10, bias=True)  
  )  
)
```




Figure 21. Testing Accuracy (80.37%)




Figure 24. Training Accuracy (99.22%)




Figure 23. Testing Loss




Figure 22. Training Loss

Part 2, Sub-Part 2 (a):

Parameters :

Name: "assignment_p2_2a_3"

Epoch = 100

Batch Size = 64

Loss function: Cross Entropy Loss

Optimizer: SGD (lr: 1e-3)

Initialization function: kaiming_normal_

Model:

```
VGG11Model(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): LeakyReLU(negative_slope=0.01, inplace=True)  
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (6): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): LeakyReLU(negative_slope=0.01, inplace=True)  
    (8): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): LeakyReLU(negative_slope=0.01, inplace=True)  
    (10): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (11): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (12): LeakyReLU(negative_slope=0.01, inplace=True)  
    (13): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (14): LeakyReLU(negative_slope=0.01, inplace=True)  
    (15): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (16): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (17): LeakyReLU(negative_slope=0.01, inplace=True)  
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (19): LeakyReLU(negative_slope=0.01, inplace=True)  
    (20): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avg_pool): AdaptiveAvgPool2d(output_size=(7, 7))  
  (classifier): Sequential(  
    (0): Linear(in_features=25088, out_features=4096, bias=True)  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): Dropout(p=0.5, inplace=False)  
    (3): Linear(in_features=4096, out_features=4096, bias=True)  
    (4): LeakyReLU(negative_slope=0.01, inplace=True)  
    (5): Dropout(p=0.5, inplace=False)  
    (6): Linear(in_features=4096, out_features=10, bias=True)  
  )  
)
```




Figure 25. Testing Accuracy (72.03%)




Figure 28. Training Accuracy (99.19%)




Figure 27. Testing Loss




Figure 26. Training Loss

Part 2, Sub-Part 2 (a) with BatchNorm:

Parameters :

Name: "assignment_p2_2a_3b"

Epoch = 100

Batch Size = 64

Loss function: Cross Entropy Loss

Optimizer: Adam (lr: 1e-3)

Initialization function: kaiming_normal_

Model:

```
VGG11Model(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): LeakyReLU(negative_slope=0.01, inplace=True)  
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (4): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (5): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (6): LeakyReLU(negative_slope=0.01, inplace=True)  
    (7): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (8): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (10): LeakyReLU(negative_slope=0.01, inplace=True)  
    (11): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (12): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (13): LeakyReLU(negative_slope=0.01, inplace=True)  
    (14): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (15): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (16): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (17): LeakyReLU(negative_slope=0.01, inplace=True)  
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (19): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (20): LeakyReLU(negative_slope=0.01, inplace=True)  
    (21): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (22): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (23): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (24): LeakyReLU(negative_slope=0.01, inplace=True)  
    (25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (26): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (27): LeakyReLU(negative_slope=0.01, inplace=True)  
    (28): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avg pool): AdaptiveAvgPool2d(output_size=(7, 7))  
  (classifier): Sequential(  
    (0): Linear(in_features=25088, out_features=4096, bias=True)  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): Dropout(p=0.5, inplace=False)  
    (3): Linear(in_features=4096, out_features=4096, bias=True)  
    (4): LeakyReLU(negative_slope=0.01, inplace=True)  
    (5): Dropout(p=0.5, inplace=False)  
    (6): Linear(in_features=4096, out_features=10, bias=True)  
  )  
)
```




Figure 30. Testing Accuracy (66.32%)




Figure 31. Training Accuracy (100%)




Figure 32. Testing Loss




Figure 29. Training Loss

Part 2, Sub-Part 2 (b):

Parameters :

Name: "assignment_p2_2b"

Epoch = 100

Batch Size = 256

Loss function: Cross Entropy Loss

Optimizer: Adam (betas:(0.5, 0.999), lr: 2e-4)

Initialization function: kaiming_normal_

Model:

```
VGG11Model(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): LeakyReLU(negative_slope=0.01, inplace=True)  
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (6): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): LeakyReLU(negative_slope=0.01, inplace=True)  
    (8): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): LeakyReLU(negative_slope=0.01, inplace=True)  
    (10): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (11): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (12): LeakyReLU(negative_slope=0.01, inplace=True)  
    (13): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (14): LeakyReLU(negative_slope=0.01, inplace=True)  
    (15): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (16): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (17): LeakyReLU(negative_slope=0.01, inplace=True)  
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (19): LeakyReLU(negative_slope=0.01, inplace=True)  
    (20): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avg_pool): AdaptiveAvgPool2d(output_size=(7, 7))  
  (classifier): Sequential(  
    (0): Linear(in_features=25088, out_features=4096, bias=True)  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): Dropout(p=0.5, inplace=False)  
    (3): Linear(in_features=4096, out_features=4096, bias=True)  
    (4): LeakyReLU(negative_slope=0.01, inplace=True)  
    (5): Dropout(p=0.5, inplace=False)  
    (6): Linear(in_features=4096, out_features=10, bias=True)  
  )  
)
```




Figure 33. Testing Accuracy (75.46%)




Figure 36. Training Accuracy (99.61)




Figure 35. Testing Loss




Figure 34. Training Loss

Part 2, Sub-Part 2 (b) with BatchNorm:

Parameters :

Name: "assignment_p2_2b_b"

Epoch = 100

Batch Size = 256

Loss function: Cross Entropy Loss

Optimizer: Adam (betas:(0.5, 0.999), lr: 2e-4)

Initialization function: kaiming_normal_

Model:

```
VGG11Model(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): LeakyReLU(negative_slope=0.01, inplace=True)  
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (4): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (5): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (6): LeakyReLU(negative_slope=0.01, inplace=True)  
    (7): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (8): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (10): LeakyReLU(negative_slope=0.01, inplace=True)  
    (11): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (12): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (13): LeakyReLU(negative_slope=0.01, inplace=True)  
    (14): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (15): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (16): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (17): LeakyReLU(negative_slope=0.01, inplace=True)  
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (19): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (20): LeakyReLU(negative_slope=0.01, inplace=True)  
    (21): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (22): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (23): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (24): LeakyReLU(negative_slope=0.01, inplace=True)  
    (25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (26): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (27): LeakyReLU(negative_slope=0.01, inplace=True)  
    (28): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avg pool): AdaptiveAvgPool2d(output_size=(7, 7))  
  (classifier): Sequential(  
    (0): Linear(in_features=25088, out_features=4096, bias=True)  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): Dropout(p=0.5, inplace=False)  
    (3): Linear(in_features=4096, out_features=4096, bias=True)  
    (4): LeakyReLU(negative_slope=0.01, inplace=True)  
    (5): Dropout(p=0.5, inplace=False)  
    (6): Linear(in_features=4096, out_features=10, bias=True)  
  )  
)
```




Figure 37. Testing Accuracy (75.52%)




Figure 40. Training Accuracy (97.31%)




Figure 39. Testing Loss




Figure 38. Training Loss

Part 2, Sub-Part 2 (c):

Parameters :

Name: "assignment_p2_2c"

Epoch = 100

Batch Size = 64

Loss function: Cross Entropy Loss

Optimizer: Adam (betas:(0.5, 0.999), lr: 2e-4)

Initialization function: xavier_uniform_

Model:

```
VGG11Model(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): LeakyReLU(negative_slope=0.01, inplace=True)  
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (6): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): LeakyReLU(negative_slope=0.01, inplace=True)  
    (8): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): LeakyReLU(negative_slope=0.01, inplace=True)  
    (10): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (11): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (12): LeakyReLU(negative_slope=0.01, inplace=True)  
    (13): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (14): LeakyReLU(negative_slope=0.01, inplace=True)  
    (15): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (16): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (17): LeakyReLU(negative_slope=0.01, inplace=True)  
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (19): LeakyReLU(negative_slope=0.01, inplace=True)  
    (20): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avg_pool): AdaptiveAvgPool2d(output_size=(7, 7))  
  (classifier): Sequential(  
    (0): Linear(in_features=25088, out_features=4096, bias=True)  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): Dropout(p=0.5, inplace=False)  
    (3): Linear(in_features=4096, out_features=4096, bias=True)  
    (4): LeakyReLU(negative_slope=0.01, inplace=True)  
    (5): Dropout(p=0.5, inplace=False)  
    (6): Linear(in_features=4096, out_features=10, bias=True)  
  )  
)
```




Figure 41. Testing Accuracy (74.1%)




Figure 44. Training Accuracy (83.15%)




Figure 43. Testing Loss




Figure 42. Training Loss

Part 2, Sub-Part 2 (c) with BatchNorm:

Parameters :

Name: "assignment_p2_2c_b"

Epoch = 100

Batch Size = 64

Loss function: Cross Entropy Loss

Optimizer: Adam (betas:(0.5, 0.999), lr: 2e-4)

Initialization function: xavier_uniform_

Model:

```
VGG11Model(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): LeakyReLU(negative_slope=0.01, inplace=True)  
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (4): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (5): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (6): LeakyReLU(negative_slope=0.01, inplace=True)  
    (7): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (8): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (10): LeakyReLU(negative_slope=0.01, inplace=True)  
    (11): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (12): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (13): LeakyReLU(negative_slope=0.01, inplace=True)  
    (14): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (15): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (16): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (17): LeakyReLU(negative_slope=0.01, inplace=True)  
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (19): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (20): LeakyReLU(negative_slope=0.01, inplace=True)  
    (21): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (22): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (23): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (24): LeakyReLU(negative_slope=0.01, inplace=True)  
    (25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (26): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (27): LeakyReLU(negative_slope=0.01, inplace=True)  
    (28): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avg pool): AdaptiveAvgPool2d(output_size=(7, 7))  
  (classifier): Sequential(  
    (0): Linear(in_features=25088, out_features=4096, bias=True)  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): Dropout(p=0.5, inplace=False)  
    (3): Linear(in_features=4096, out_features=4096, bias=True)  
    (4): LeakyReLU(negative_slope=0.01, inplace=True)  
    (5): Dropout(p=0.5, inplace=False)  
    (6): Linear(in_features=4096, out_features=10, bias=True)  
  )  
)
```




Figure 45. Testing Accuracy (82.05%)




Figure 48. Training Accuracy (98.8%)




Figure 47. Testing Loss




Figure 46. Training Loss

Part 2, Sub-Part 2 (d):

Parameters :

Name: "assignment_p2_2d"

Epoch = 100

Batch Size = 64

Loss function: Cross Entropy Loss

Optimizer: Adam (betas:(0.5, 0.999), lr: 2e-4)

Initialization function: kaiming_normal_

Model:

```
VGG11Model(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (4): LeakyReLU(negative_slope=0.01, inplace=True)  
    (5): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (6): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): LeakyReLU(negative_slope=0.01, inplace=True)  
    (8): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): LeakyReLU(negative_slope=0.01, inplace=True)  
    (10): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (11): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (12): LeakyReLU(negative_slope=0.01, inplace=True)  
    (13): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (14): LeakyReLU(negative_slope=0.01, inplace=True)  
    (15): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (16): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (17): LeakyReLU(negative_slope=0.01, inplace=True)  
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (19): LeakyReLU(negative_slope=0.01, inplace=True)  
    (20): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avg_pool): AdaptiveAvgPool2d(output_size=(7, 7))  
  (classifier): Sequential(  
    (0): Linear(in_features=25088, out_features=4096, bias=True)  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): Linear(in_features=4096, out_features=4096, bias=True)  
    (3): LeakyReLU(negative_slope=0.01, inplace=True)  
    (4): Linear(in_features=4096, out_features=10, bias=True)  
  )  
)
```




Figure 49. Testing Accuracy (70.4%)




Figure 51. Testing Loss




Figure 50. Training Loss

Figure 52. Training Accuracy (98.86%)

Part 2, Sub-Part 2 (d) with BatchNorm:

Parameters :

Name: "assignment_p2_2d_b"

Epoch = 100

Batch Size = 64

Loss function: Cross Entropy Loss

Optimizer: Adam (betas:(0.5, 0.999), lr: 2e-4)

Initialization function: kaiming_normal_

Model:

```
VGG11Model(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): LeakyReLU(negative_slope=0.01, inplace=True)  
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (4): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (5): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (6): LeakyReLU(negative_slope=0.01, inplace=True)  
    (7): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (8): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (10): LeakyReLU(negative_slope=0.01, inplace=True)  
    (11): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (12): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (13): LeakyReLU(negative_slope=0.01, inplace=True)  
    (14): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (15): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (16): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (17): LeakyReLU(negative_slope=0.01, inplace=True)  
    (18): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (19): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (20): LeakyReLU(negative_slope=0.01, inplace=True)  
    (21): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (22): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (23): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (24): LeakyReLU(negative_slope=0.01, inplace=True)  
    (25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (26): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (27): LeakyReLU(negative_slope=0.01, inplace=True)  
    (28): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avg pool): AdaptiveAvgPool2d(output_size=(7, 7))  
  (classifier): Sequential(  
    (0): Linear(in_features=25088, out_features=4096, bias=True)  
    (1): LeakyReLU(negative_slope=0.01, inplace=True)  
    (2): Linear(in_features=4096, out_features=4096, bias=True)  
    (3): LeakyReLU(negative_slope=0.01, inplace=True)  
    (4): Linear(in_features=4096, out_features=10, bias=True)  
  )  
)
```




Figure 53. Training Accuracy (82.06%)




Figure 55. Training Accuracy (99.66%)




Figure 56. Testing Loss




Figure 54. Training Loss